

## PySpark\_WorkingWithCSV\_HandsOn\_3Sep

SaiPrabath Chowdary S

```
✓ 37s [1] # https://codeshare.io/pAe0RV  
from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
✓ 45s [2] ! pip install pyspark
```

collecting pyspark  
Downloading pyspark-3.5.2.tar.gz (317.3 MB) 317.3/317.3 MB 4.1 MB/s eta 0:00:00  
Preparing metadata (setup.py) ... done  
Requirement already satisfied: pyd==0.10.9.7 in /usr/local/lib/python3.10/dist-packages (from pyspark) (0.10.9.7)  
Building wheels for collected packages: pyspark  
Building wheel for pyspark (setup.py) ... done  
Created wheel for pyspark: filename=pyspark-3.5.2-py2.py3-none-any.whl size=317812365 sha256=b3f121c3595d4f009275ab679fe29198a4f0e64c58d94b6e83614d570a24701e  
Stored in directory: /root/.cache/pip/wheels/34/34/bd/03944534c44b677cd5859f248090daa9fb27b3c8f8e5f49574  
Successfully built pyspark  
Installing collected packages: pyspark  
Successfully installed pyspark-3.5.2

```
▶ # step 1  
import pandas as pd  
from datetime import datetime  
  
data = {  
    "TransactionID": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],  
    "CustomerID": [101, 102, 103, 101, 104, 102, 103, 104, 101, 105],  
    "ProductID": [501, 502, 501, 503, 504, 502, 503, 504, 501, 505],  
    "Quantity": [2, 1, 4, 3, 1, 2, 5, 1, 2, 1],  
    "Price": [150.0, 250.0, 150.0, 300.0, 450.0, 250.0, 300.0, 450.0, 150.0, 550.0],  
    "Date": [  
        datetime(2024, 9, 1),  
        datetime(2024, 9, 1),  
        datetime(2024, 9, 2),  
        datetime(2024, 9, 2),  
        datetime(2024, 9, 3),  
        datetime(2024, 9, 3),  
        datetime(2024, 9, 4),  
        datetime(2024, 9, 4),  
        datetime(2024, 9, 5),  
        datetime(2024, 9, 5)  
    ]  
}  
  
# DataFrame  
df = pd.DataFrame(data)  
  
# Save the DataFrame to a CSV file  
df.to_csv('/content/drive/MyDrive/DataEngineering/sales_data.csv', index=False)  
  
print("Sample sales dataset has been created and saved as 'sales_data.csv'.")
```

Sample sales dataset has been created and saved as 'sales\_data.csv'.

✓  
1s

# step 2

from pyspark.sql import SparkSession  
  
spark = SparkSession.builder.appName("Sales Dataset Analysis").getOrCreate()  
  
sales\_df = spark.read.csv("/content/drive/MyDrive/DataEngineering/sales\_data.csv", header=True, inferSchema=True)  
  
sales\_df.show(4) #verify

⇌

TransactionID	CustomerID	ProductID	Quantity	Price	Date
1	101	501	2	150.0	2024-09-01
2	102	502	1	250.0	2024-09-01
3	103	501	4	150.0	2024-09-02
4	101	503	3	300.0	2024-09-02

only showing top 4 rows

✓  
2s

[6] # step 3

# printing schema  
sales\_df.printSchema()  
  
# first 5 rows  
sales\_df.show(5)  
  
# summary statistics for numeric columns  
sales\_df.describe(["Quantity", "Price"]).show()

⇌

```
root
|-- TransactionID: integer (nullable = true)
|-- CustomerID: integer (nullable = true)
|-- ProductID: integer (nullable = true)
|-- Quantity: integer (nullable = true)
|-- Price: double (nullable = true)
|-- Date: date (nullable = true)
```

TransactionID	CustomerID	ProductID	Quantity	Price	Date
1	101	501	2	150.0	2024-09-01
2	102	502	1	250.0	2024-09-01
3	103	501	4	150.0	2024-09-02
4	101	503	3	300.0	2024-09-02
5	104	504	1	450.0	2024-09-03

only showing top 5 rows

summary	Quantity	Price
count	10	10
mean	2.2	300.0
stddev	1.398411797560202	141.4213562373095
min	1	150.0
max	5	550.0

✓ [19] # step 4

```
# 1. Calculate the Total Sales Value for Each Transaction
sales_df = sales_df.withColumn("TotalSales", sales_df["Quantity"] * sales_df["Price"])
sales_df.show()
```



TransactionID	CustomerID	ProductID	Quantity	Price	Date	TotalSales
1	101	501	2	150.0	2024-09-01	300.0
2	102	502	1	250.0	2024-09-01	250.0
3	103	501	4	150.0	2024-09-02	600.0
4	101	503	3	300.0	2024-09-02	900.0
5	104	504	1	450.0	2024-09-03	450.0
6	102	502	2	250.0	2024-09-03	500.0
7	103	503	5	300.0	2024-09-04	1500.0
8	104	504	1	450.0	2024-09-04	450.0
9	101	501	2	150.0	2024-09-05	300.0
10	105	505	1	550.0	2024-09-05	550.0

✓ [20] # 2. Group By ProductID and Calculate Total Sales Per Product

```
total_sales_by_product = sales_df.groupBy("ProductID").sum("TotalSales").withColumnRenamed("sum(TotalSales)", "TotalSales")\
    .orderBy("sum(TotalSales)", ascending=False)
total_sales_by_product.show()
```



ProductID	TotalSales
503	2400.0
501	1200.0
504	900.0
502	750.0
505	550.0

[21] # 3. Identify the Top-Selling Product

```
top_selling_product = total_sales_by_product.limit(1)
top_selling_product.show()
```



ProductID	TotalSales
503	2400.0



# 4. Calculate the Total Sales by Date

```
total_sales_by_date = sales_df.groupBy("Date").sum('TotalSales').orderBy("Date")
total_sales_by_date.show()
```



Date	sum(TotalSales)
2024-09-01	550.0
2024-09-02	1500.0
2024-09-03	950.0
2024-09-04	1950.0
2024-09-05	850.0

✓  
0s



#### # 5.Filter High-Value Transactions

```
high_value_transactions = sales_df.filter(sales_df["TotalSales"] > 500)
high_value_transactions.show()
```



TransactionID	CustomerID	ProductID	Quantity	Price	Date	TotalSales
3	103	501	4	150.0	2024-09-02	600.0
4	101	503	3	300.0	2024-09-02	900.0
7	103	503	5	300.0	2024-09-04	1500.0
10	105	505	1	550.0	2024-09-05	550.0

✓  
1s

#### [24] # 6.Identify Repeat Customers

```
repeat_customers = sales_df.groupBy("CustomerID").count().filter("count > 1").orderBy("count", ascending=False)
repeat_customers.show()
```



CustomerID	count
101	3
103	2
102	2
104	2

✓  
0s



#### # 7.Calculate the Average Sale Price Per Product

```
average_sale_price = sales_df.groupBy("ProductID").avg('Price').withColumnRenamed("avg(Price)", "AveragePrice").orderBy("AveragePrice", ascending=False)
average_sale_price.show()
```



ProductID	AveragePrice
505	550.0
504	450.0
503	300.0
502	250.0
501	150.0

